

WEEK 2

RDBMS ARCHITECTURE

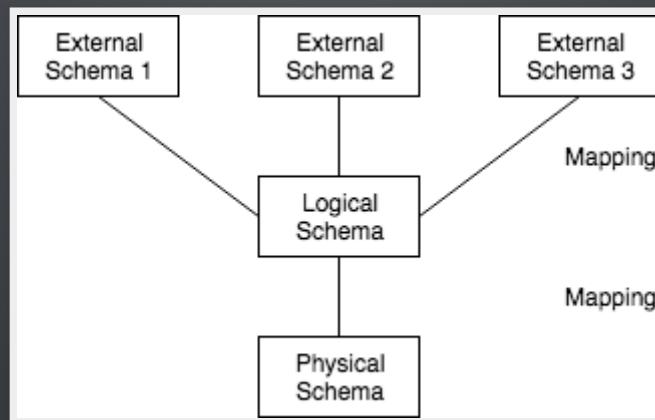
NORMALIZATION

MORE SQL

DATABASE ARCHITECTURE

3 SCHEMA ARCHITECTURE

- External / User Views
- Conceptual / Logical
- Internal / Physical



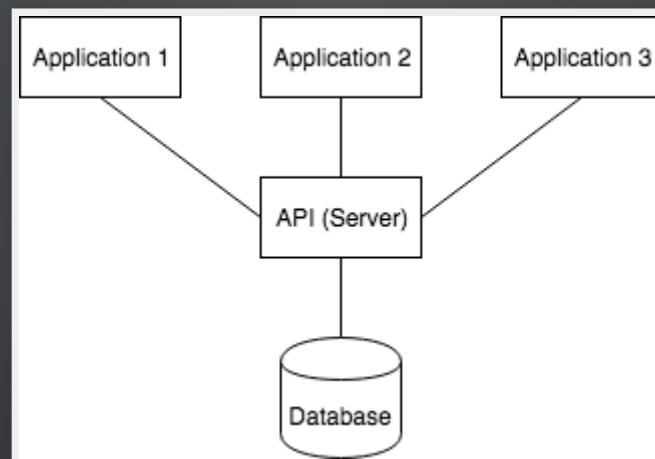
SCHEMA LANGUAGES

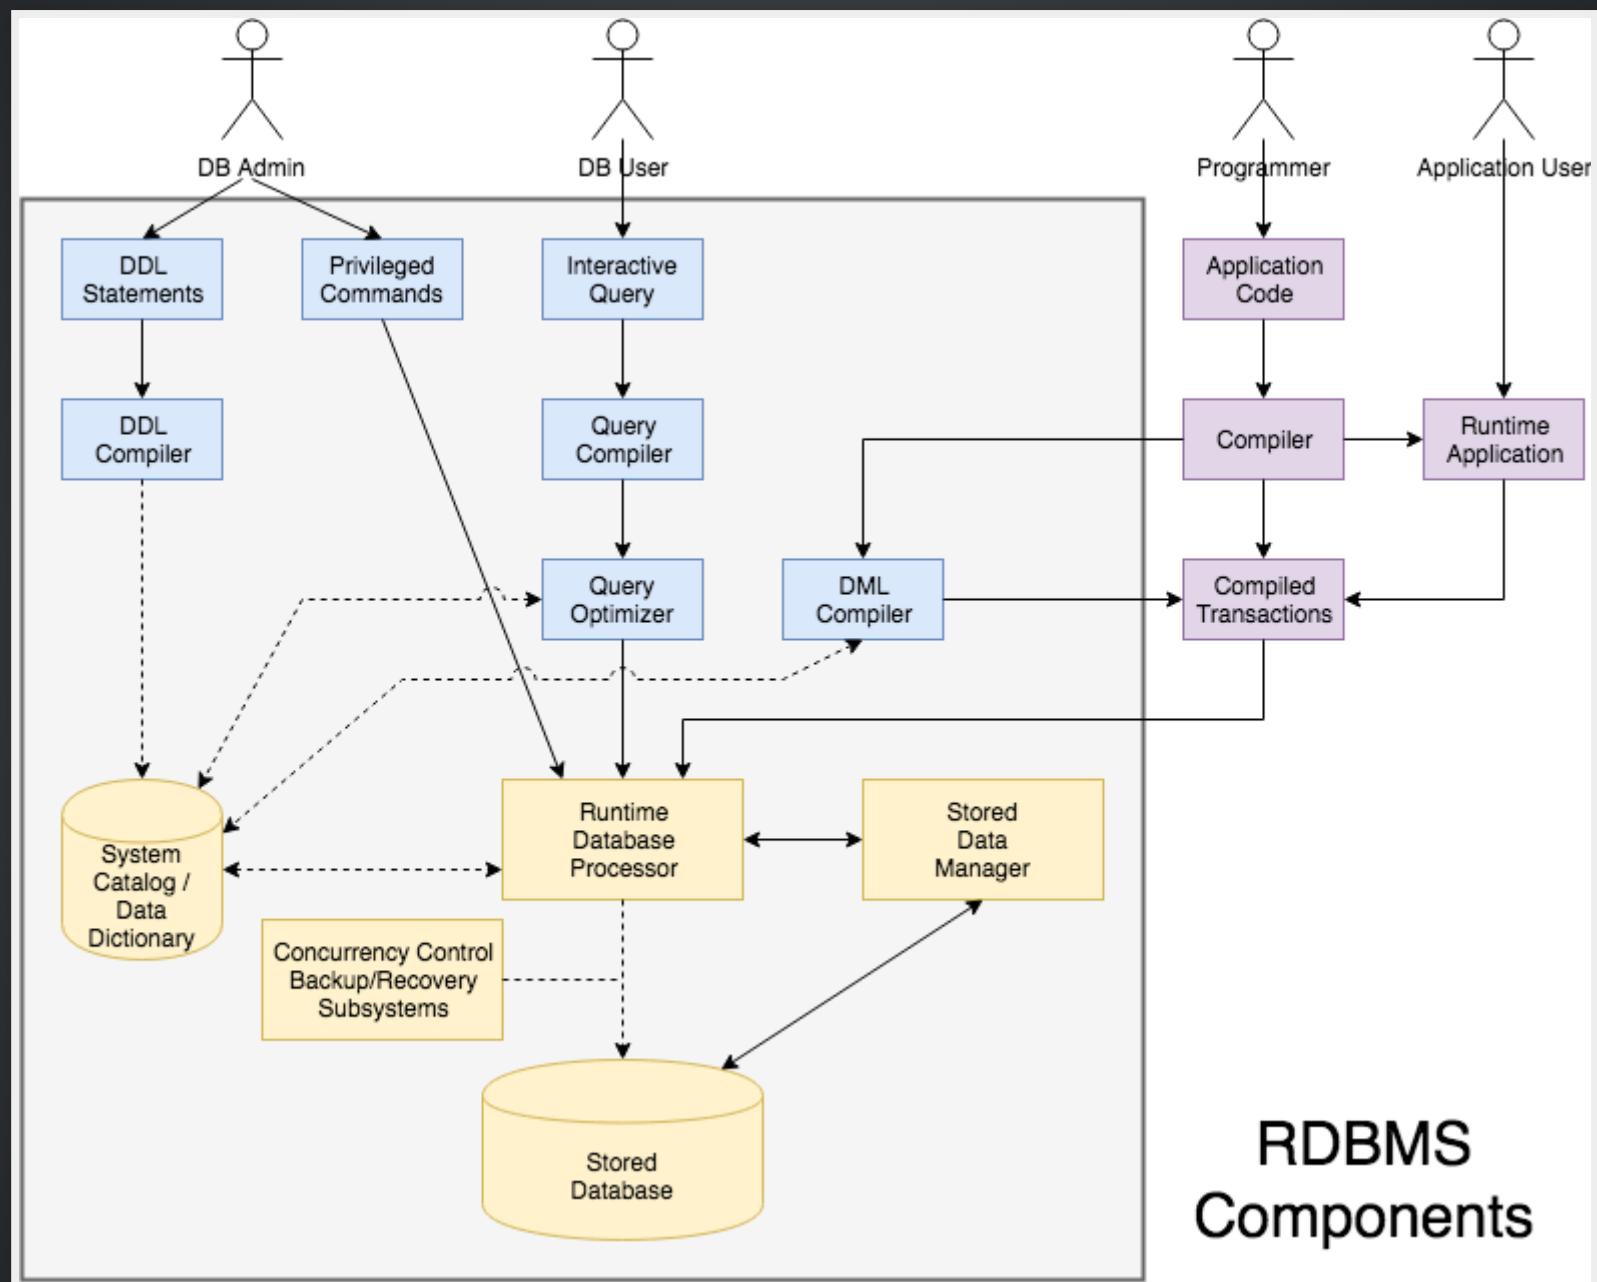
- VDL: view definition
- DML: data manipulation
- DDL: data definition
- SDL: storage definition

-
- All rolled into SQL
 - Except SDL

3 TIER ARCHITECTURE

- Client Application(s)
- API (Server)
- Database





RDBMS UTILITIES

- Loading
- Backup
- Monitoring
- Reorganization
 - Cleans up mess from deleted rows

SCHEMA NORMALIZATION

RULES OF THUMB

1. Avoid Ambiguous Attributes
2. Avoid Redundant Data
3. Avoid NULLs
4. Avoid Spurious JOINS

1. AVOID AMBIGUOUS ATTRIBUTES

- Semantics of the attributes should be clear
- Design tables that are easy to explain
- Don't mix attributes from multiple entity types in a single relation

OBJECTS

| ObjectID | Title | Type | DateCreated | ArtistID | ArtistName | ArtistDOB |
|----------|---------|----------|-------------|----------|----------------------|------------|
| 1 | Irises | Painting | 1889 | 1 | Van Gogh, Vincent | 1853-03-30 |
| 2 | Spring | Painting | 1881 | 2 | Manet, Édouard | 1832-01-23 |
| 3 | Olympia | Painting | 1862 | 2 | Manet, Édouard | 1832-01-23 |

EXHIBITION_LOANS

| ObjID | ExName | ExStart | ExEnd | LoanStart | LoanEnd | Museum | Mgr | Address | Phone |
|-------|---------------|----------------|----------------|----------------|----------------|------------------|---------|------------------|-------------------------------|
| 1 | Impress Me | 2018- 01-01 | 2018- 06-01 | 2017-12- 01 | 2018-07- 07 | The Getty | PersonX | LA, CA, USA | (310) 440- 7300 |
| 2 | Impress Me | 2018- 01-01 | 2018- 06-01 | 2017-11- 09 | 2018-08- 01 | The Getty | PersonY | LA CA, USA | (310) 440- 7300 |
| 3 | Impress Me | 2018- 01-01 | 2018- 06-01 | 2017-10- 11 | 2018-06- 22 | Musée d'Orsay | PersonZ | Paris, France | +33 (0)1 40 49 48 14 |

2. AVOID REDUNDANT DATA

- Repeated values wastes storage space
- Also leads to anomalies
 - insertion anomalies
 - deletion anomalies
 - modification anomalies

INSERTION ANOMALIES

OBJECTS

| ObjectID | Title | Type | DateCreated | ArtistID | ArtistName | ArtistDOB |
|----------|---------|----------|-------------|----------|----------------------|------------|
| 1 | Irises | Painting | 1889 | 1 | Van Gogh, Vincent | 1853-03-30 |
| 2 | Spring | Painting | 1881 | 2 | Manet, Édouard | 1832-01-23 |
| 3 | Olympia | Painting | 1862 | 2 | Manet, Édouard | 1832-01-23 |

- Inserting an object requires artist data (or NULLs)
 - Prone to inconsistency
- How do you insert an artist w/o an object?
 - Use NULLs
 - After an object is inserted w/ that artist?

DELETION ANOMALIES

OBJECTS

| ObjectID | Title | Type | DateCreated | ArtistID | ArtistName | ArtistDOB |
|----------|---------|----------|-------------|----------|----------------------|------------|
| 1 | Irises | Painting | 1889 | 1 | Van Gogh, Vincent | 1853-03-30 |
| 2 | Spring | Painting | 1881 | 2 | Manet, Édouard | 1832-01-23 |
| 3 | Olympia | Painting | 1862 | 2 | Manet, Édouard | 1832-01-23 |

- What if you delete the last object referring to a particular artist?
 - Artist is lost from the database

MODIFICATION ANOMALIES

OBJECTS

| ObjectID | Title | Type | DateCreated | ArtistID | ArtistName | ArtistDOB |
|----------|---------|----------|-------------|----------|----------------------|------------|
| 1 | Irises | Painting | 1889 | 1 | Van Gogh, Vincent | 1853-03-30 |
| 2 | Spring | Painting | 1881 | 2 | Manet, Édouard | 1832-01-23 |
| 3 | Olympia | Painting | 1862 | 2 | Manet, Édouard | 1832-01-23 |

- What if we update the Manager name for a museum?
 - We must update all of the rows with that museum
 - Leads to inconsistency

3. AVOID NULLS

- NULLs also waste space
- Can lead to ambiguity when performing:
 - JOINs
 - aggregate operations (count, sum, max)
 - comparisons (WHERE $x > y$)
- NULLs can have multiple meanings:
 - N/A, unknown, missing

4. AVOID SPURIOUS JOINS

- Design tables that can be joined on Primary Key-Foreign Key pairs
- Joining on non-prime attributes leads to spurious tuples

OBJECTS

| <u>Title</u> | <u>Type</u> | <u>DateCreated</u> | <u>ArtistName</u> | <u>ArtistDOB</u> |
|--------------|-------------|--------------------|----------------------|------------------|
| Irises | Painting | 1889 | Van Gogh, Vincent | 1853-03-30 |
| Spring | Painting | 1881 | Manet, Édouard | 1832-01-23 |
| Olympia | Painting | 1862 | Manet, Édouard | 1832-01-23 |
| Spring | Sculpture | 2001 | Dude, Some | NULL |

EXHIBITION_LOANS

| <u>Title</u> | <u>ExName</u> | <u>ExStart</u> | <u>ExEnd</u> | <u>LoanStart</u> | <u>LoanEnd</u> | <u>Museum</u> | <u>Mgr</u> | <u>Address</u> |
|--------------|---------------|----------------|----------------|------------------|----------------|------------------|------------|------------------|
| Irises | Impress Me | 2018- 01-01 | 2018- 06-01 | 2017-12- 01 | 2018-07- 07 | The Getty | PersonX | LA, CA, USA |
| Spring | Impress Me | 2018- 01-01 | 2018- 06-01 | 2017-11- 09 | 2018-08- 01 | The Getty | PersonX | LA CA, USA |
| Olympia | Impress | 2018- 01-01 | 2018- 06-01 | 2017-10- 11 | 2018-06- 22 | Musée d'Orsay | PersonZ | Paris, France |

SPURIOUS JOIN

```
SELECT o.Title, o.ArtistName, e.Museum  
FROM Objects o, Exhibition_Loans e  
WHERE e.ExName="Impress Me"  
AND e.Title = o.Title;
```

| Title | ArtistName | Museum |
|---------|-------------------|---------------|
| Irises | Van Gogh, Vincent | The Getty |
| Spring | Manet, Édouard | The Getty |
| Olympia | Manet, Édouard | Musée d'Orsay |
| Spring | Dude, Some | The Getty |

NORMALIZATION

NORMALIZATION BASICS

- A formal method of analysis to minimize redundancy and anomalies
- Has multiple forms (levels), higher is better
- Based on primary keys and functional dependencies
- Normalizing tables individually is not sufficient
 - Taken together, the tables must also ensure nonadditive joins (no spurious tuples)

FUNCTIONAL DEPENDENCY (FD)

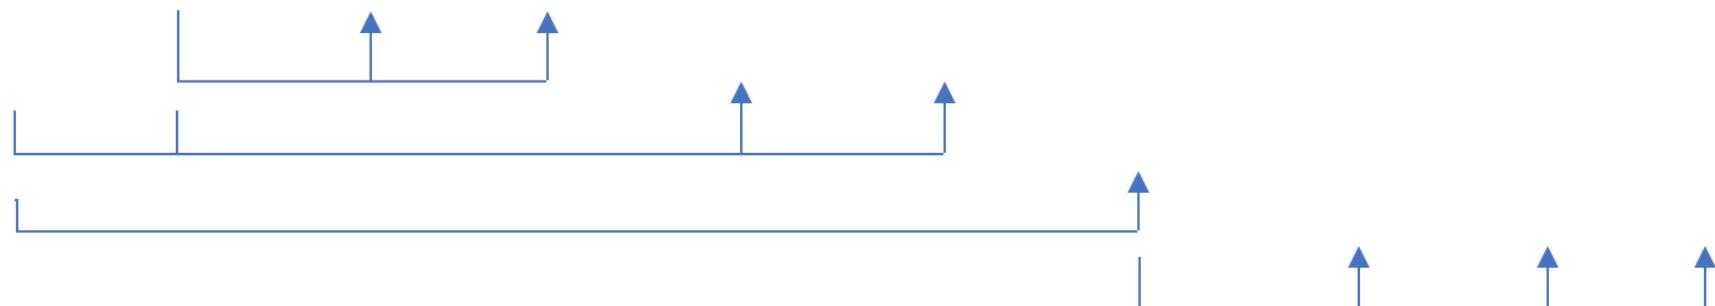
- A constraint between 2 sets of attributes
- Denoted by $X \rightarrow Y$
 - X and Y can each have more than one attribute
 - Means that the values of Y are determined by the values of X

For any 2 tuples in the database that have the same set of X values, they must also have the same set of Y values.

FUNCTIONAL DEPENDENCIES

EXHIBITION_LOANS

| ObjID | ExName | ExStart | ExEnd | LoanStart | LoanEnd | Museum | Mgr | Address | Phone |
|-------|------------|------------|------------|------------|------------|---------------|---------|---------------|----------------------|
| 1 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-12-01 | 2018-07-07 | The Getty | PersonX | LA, CA, USA | (310) 440-7300 |
| 2 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-11-09 | 2018-08-01 | The Getty | PersonX | LA CA, USA | (310) 440-7300 |
| 3 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-10-11 | 2018-06-22 | Musée d'Orsay | PersonZ | Paris, France | +33 (0)1 40 49 48 14 |



FIRST NORMAL FORM (1NF)

*Values of any attribute must be atomic
and singular*

MOVIES

| Title | Year | Genres | |
|--------------------|------|---------------------------|---|
| The Shape of Water | 2017 | Fantasy, Romance | X |
| Get Out | 2017 | Horror, Mystery, Thriller | |

TECHNIQUES FOR 1NF

1. Move the violating attribute to another table
 - e.g. Movies, Genres, HasGenre
 - Decomposition
2. Repeat the row for each value in the attribute
 - Creates redundancy
3. Add attributes for each possible value
 - e.g. Genre1, Genre2, Genre3, etc.
 - Must know maximum number of values
 - Creates NULLs

2ND NORMAL FORM (2NF)

Every nonprime attribute in a table is fully functionally dependent (FFD) on the primary key

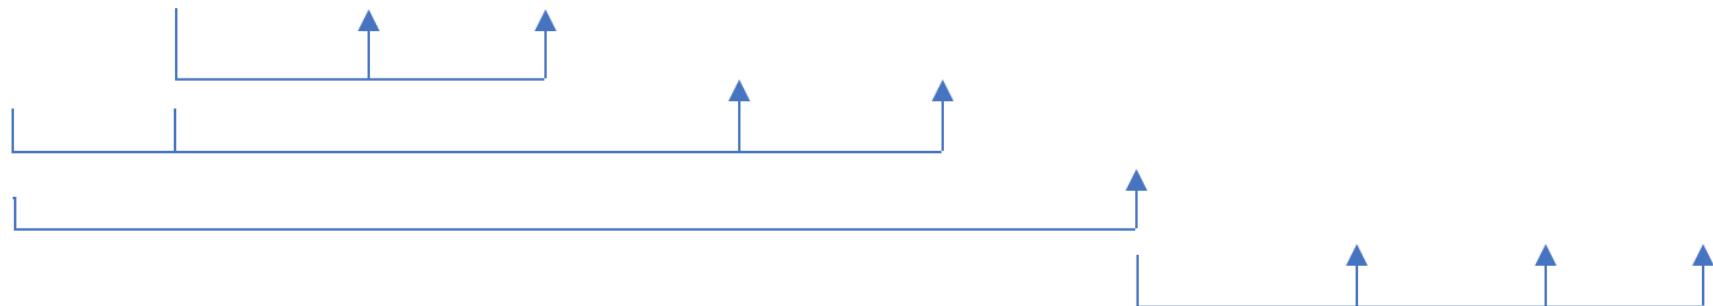
- Nonprime = not part of the primary key (PK)
- FFD = removal of any part of the PK erodes the FD
- Only applies to tables with composite PKs

TECHNIQUE FOR 2NF

- Decomposition: new table for each partial key

EXHIBITION_LOANS

| ObjID | ExName | ExStart | ExEnd | LoanStart | LoanEnd | Museum | Mgr | Address | Phone |
|-------|------------|------------|------------|------------|------------|---------------|---------|---------------|----------------------|
| 1 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-12-01 | 2018-07-07 | The Getty | PersonX | LA, CA, USA | (310) 440-7300 |
| 2 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-11-09 | 2018-08-01 | The Getty | PersonX | LA CA, USA | (310) 440-7300 |
| 3 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-10-11 | 2018-06-22 | Musée d'Orsay | PersonZ | Paris, France | +33 (0)1 40 49 48 14 |



3RD NORMAL FORM (3NF)

No nonprime attribute is transitively dependent on the primary key

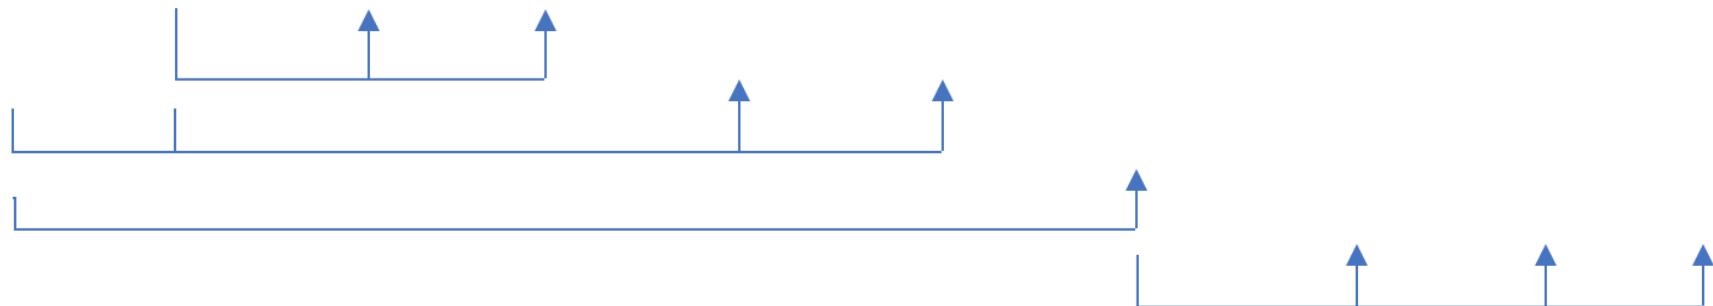
- i.e. All attributes should be FFD on the PK, not on another attribute

TECHNIQUE FOR 3NF

- Decompose: new table for transitive dependencies

EXHIBITION_LOANS

| ObjID | ExName | ExStart | ExEnd | LoanStart | LoanEnd | Museum | Mgr | Address | Phone |
|-------|------------|------------|------------|------------|------------|---------------|---------|---------------|----------------------|
| 1 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-12-01 | 2018-07-07 | The Getty | PersonX | LA, CA, USA | (310) 440-7300 |
| 2 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-11-09 | 2018-08-01 | The Getty | PersonX | LA CA, USA | (310) 440-7300 |
| 3 | Impress Me | 2018-01-01 | 2018-06-01 | 2017-10-11 | 2018-06-22 | Musée d'Orsay | PersonZ | Paris, France | +33 (0)1 40 49 48 14 |



BOYCE-CODD NORMAL FORM (BCNF)

Essentially the same as 3NF, with small addition:

A nonprime attribute cannot functionally determine a prime attribute

- Any table that is BCNF, is also 3NF
- Most tables in 3NF are also BCNF

4TH NORMAL FORM (4NF)

A table cannot have a multivalued dependency (MVD)

- MVD: occurs when 2 independent 1:N relationships exist in the same table
 - Results in repeated values
 - Often occurs in all-key tables

TECHNIQUE FOR 4NF

- Decompose each dependency into a new table

| <u>Band</u> | <u>Album</u> | <u>Musician</u> |
|-------------------------|--------------------|------------------|
| Queens of the Stone Age | Like Clockwork | Josh Homme |
| Queens of the Stone Age | Songs for the Deaf | Josh Homme |
| Queens of the Stone Age | Songs for the Deaf | Troy Van Leeuwen |
| Queens of the Stone Age | Songs for the Deaf | Dave Grohl |
| Nirvana | Nevermind | Dave Grohl |
| Nirvana | Nevermind | Kurt Cobain |
| Nirvana | Nevermind | Chris Novaselic |

5TH NORMAL FORM?

- Dont' bother
 - For join dependencies among 3 or more tables
- In practice, most DBAs aim for 3NF, which is critical
- BCNF, 4NF, and 5NF are subtle and rarely needed

MORE SQL

ON DELETE

What should happen if a foreign key is deleted?

```
CREATE TABLE chapter(
    number  INTEGER,
    title   VARCHAR(256),
    bookID  INTEGER REFERENCES book ON DELETE CASCADE
);
```

Options include: SET DEFAULT, SET NULL, CASCADE,
RESTRICT, NO ACTION

ON UPDATE

What happens if the foreign key is changed?

e.g. FK is an ISBN, changed from 10 to 13 digits

```
CREATE TABLE chapter(
    number  INTEGER,
    title   VARCHAR(256),
    isbn    VARCHAR(15) REFERENCES book.isbn
            ON DELETE CASCADE
            ON UPDATE CASCADE
);
```

Same options as ON DELETE

ALTER TABLE

Want to rename a table?

```
ALTER TABLE actors RENAME TO thesbians;
```

What if you want to add or remove a column?

```
ALTER TABLE chapter ADD COLUMN pageCount INTEGER;
```

```
ALTER TABLE chapter DROP COLUMN pageCount;
```

DROP COLUMN not supported by some DBs

MORE ABOUT SELECT

- DISTINCT
- LIKE
- LIMIT &
OFFSET

DISTINCT

- Imagine a DB with 3 tables:
 - musician: id, name
 - album: id, title, year
 - played_on: musicianID, albumID, role
- How do you get all the roles a musician has ever performed without repeats?

```
SELECT DISTINCT musician.name, played_on.role  
FROM musician, played_on  
WHERE played_on.musicianID=musician.id;
```

LIKE

- Fuzzy matching for text with regular expressions
- Wildcard Characters
 - % represents zero, one, or more characters
 - _ represents a single character
- e.g. Find books with "information" in the title

```
SELECT * FROM book WHERE title LIKE '%information%';
```

LIMIT

Too many results?

LIMIT them!

```
SELECT * FROM books WHERE title LIKE '%the%' LIMIT 20;
```

OFFSET

What if you want to page through the results?

Use LIMIT and OFFSET together

```
SELECT * FROM books WHERE title LIKE '%the%' LIMIT 20 OFFSET 20;
```

NESTED QUERIES

- Sometimes you need to perform one query and use the results in another
- You can do that in one statement with nested queries

```
SELECT title
FROM album, band
WHERE album.bandid=band.id
AND band.name="Metallica"
AND album.id IN (
    SELECT id
    FROM album
    ORDER BY earnings DESC
    LIMIT 100
);
```

AGGREGATE FUNCTIONS

- Standard SQL has aggregate functions that can be called within a query.
- Each RDBMS also supports it's own functions
- SQLite3 supports these functions:
 - `min(X)`, `max(X)`, `avg(X)`, `sum(X)`, `total(X)`
 - `count(*)`, `count(X)`
 - `group_concat(X, Y)`

```
SELECT count(title) FROM films;
```

```
SELECT max(earnings), avg(earnings) FROM films;
```

OUTER JOINS

- What happens when you make a query using a FK that can have NULLs?
 - If you use a NATURAL INNER JOIN (aka equijoin), the rows with NULLs will not be included

```
SELECT bookt.title, author.name  
FROM book, author  
WHERE book.authorid=author.id;
```

- If you want to include rows with NULLs you have to do an OUTER JOIN

OUTER JOINS

- To keep all the elements from the first table use a
LEFT OUTER JOIN

```
SELECT book.title, author.name  
FROM book  
LEFT OUTER JOIN ON book.authorid=author.id;
```

| Title | Name |
|-------------------|------------------|
| Ancillary Justice | Ann Leckie |
| Arabian Nights | |
| Night Circus | Erin Morgenstern |

OUTER JOINS

- To keep all the elements from the second table use a
RIGHT OUTER JOIN

```
SELECT book.title, author.name
FROM book
RIGHT OUTER JOIN ON book.authorid=author.id;
```

| Title | Name |
|-------------------|------------------|
| Ancillary Justice | Ann Leckie |
| | Joshua Gomez |
| Night Circus | Erin Morgenstern |

OUTER JOINS

- To keep all the elements from both tables use a FULL OUTER JOIN

```
SELECT book.title, author.name  
FROM book  
FULL OUTER JOIN ON book.authorid=author.id;
```

| Title | Name |
|-------------------|------------------|
| Ancillary Justice | Ann Leckie |
| Arabian Nights | |
| Night Circus | Erin Morgenstern |
| | Joshua Gomez |

GROUP BY

- You can arrange your results in groups
- This works well with aggregate functions

```
SELECT count(id), genre
FROM albums
GROUP BY genre
ORDER BY count(id) DESC;
```

HAVING

- What if you want to add a condition based on the aggregate function?
- Use the HAVING clause

```
SELECT count(id), genre
FROM albums
GROUP BY genre
HAVING count(id) > 1000
ORDER BY count(id) DESC;
```

EOL

On to the exercise!

After a break